

AI LITERACY FOR EVERYONE FOUNDATION


AI 101

How Did We Get Here?
Where Are We Going?

May 2025

Sven Olensky

svenolensky@AiLiteracyForEveryone.org



Contents

Note about this document.....	3
About the Presenter	3
Introduction: The Unseen History of AI	4
AI is an evolution of personal computing	4
Key Takeaways	4
AI in Historical Context: The Building Blocks of Intelligence.....	4
The Beginnings: Turing Machine	4
Expert Systems and Rule-Based AI	5
Shift to Learning Models	5
Key Takeaways	6
How AI Works: A High-Level Primer	6
Key Takeaways	7
The AI Boom: AI in Everyday Life.....	8
Key Takeaways	11
AI's Role in Society: Opportunities & Challenges	11
Opportunities and Benefits	11
Challenges and Concerns.....	12
Key Takeaways	14
Looking Ahead: Where AI is Taking Us	14
Quantum Computing: The Next Frontier of Speed and Intelligence	17
Key Takeaways	19
The Future of Learning: AI as a Cognitive Amplifier.....	19
Optimizing Learning through Pareto's 80/20 approach	20
Key Takeaways	21
AI Trust & Verification: How to Evaluate AI-Generated Information.....	22
Objective	22
Why AI Can't Be Trusted Blindly.....	22
How to Verify AI Outputs	22
Spotting AI Bias and Misinformation	23
Real-World Consequences of Blindly Trusting AI	23
Applying 'Trust But Verify' to AI-Assisted Learning	23
Some Ideas to Experiment with	24

Key Takeaways	24
The Future Is Yours	24
Key Takeaways	25
References	26

Note about this document

This whitepaper was created using ChatGPT to demonstrate the current capabilities of AI-assisted content generation. While formatting, layout, and visuals (many of which are also AI-generated) were added for clarity and engagement, the core narrative remains nearly 100% AI-generated—based on curated inputs and editorial direction.

About the Presenter

Sven D. Olensky, CBSE, AWS-AIP, CAISS, CCZT, AWS-SCS, CCSP, ISSAP, CISSP, is a Senior Principal Security Architect with over 30 years of experience in technology and cybersecurity.

At Cox Enterprises, Sven leads the development of enterprise-wide security architecture strategies—aligning cybersecurity with business goals and an evolving threat landscape. His current work spans security governance, multi-cloud frameworks, automated compliance guardrails, and data protection architecture. He collaborates closely with platform teams, solution architects, and business leaders to embed security into technology roadmaps, while mentoring others to raise the organization’s overall security maturity.

Previously, during nearly a decade at the Federal Reserve, he spearheaded efforts in AI Security, Zero Trust, and cloud architecture—safeguarding systems that process more than \$5 trillion daily. He authored the Federal Reserve System’s practical AI Security Standard and contributed to multiple Cloud Security Alliance (CSA) standards on AI and Zero Trust.

Sven’s career includes roles at Newell Brands, Home Depot, Cox Communications, Dell SecureWorks, and others, where he led secure architecture design, implemented enterprise controls, and drove regulatory compliance initiatives. He has shaped industry certifications by contributing to (ISC)² exam development for the CISSP and CCSP. As founder of GANSEC, he brings enterprise-grade cybersecurity guidance to small and mid-sized businesses.

His passion for technology began with *Tron* and *WarGames*, and deepened when he taught himself to code on a Commodore 64. From building PCs to exploring emerging tech, Sven has spent nearly four decades engineering, problem-solving, and teaching—driven by a lifelong commitment to curiosity and innovation.

Introduction: The Unseen History of AI

Artificial Intelligence (AI) may seem like a cutting-edge trend, but its roots stretch back over half a century. The **myth that AI is a recent innovation** is just that: a myth. Researchers have pursued the dream of intelligent machines since the mid-20th century. The term “*artificial intelligence*” itself was coined in **1956** at a famed workshop at Dartmouth College, marking the birth of AI as an academic field. Even earlier, in 1950, **Alan Turing** pondered machine intelligence in his seminal paper on the *Imitation Game* (known today as the Turing Test). These early explorations laid the groundwork for what would become modern AI.

AI is an evolution of personal computing

AI is not an alien newcomer. Just as personal computers revolutionized work and play, AI is a natural next step—enabling computers to *learn* and *adapt* rather than just follow fixed instructions. In the 1980s, for instance, a supercomputer like the Cray-2 could do in one day all the computation NASA had done in the year 1970. Today, your smartphone has comparable power. As computing power grew exponentially (often following Moore’s Law), researchers kept the flame of AI alive through decades of ups and downs. What feels like an overnight AI “boom” is the result of **decades of iterative progress**, from early experiments to the powerful learning systems now in the palms of our hands.

In this presentation and workshop, we’ll journey from **AI’s unseen history** to the present and beyond. We’ll see how AI evolved alongside computers, starting as lofty theory in lab memos and becoming everyday reality in our apps and gadgets. By understanding where AI came from, we can better appreciate where it’s headed and how it will continue to shape our world.

Key Takeaways

- AI has been evolving for over **70 years**—what feels new is decades of research reaching practical applications.
- The rise of AI is **tied to computing progress**—from mainframes to personal computers to AI-powered devices.
- **Understanding AI’s past** helps us grasp its future—AI is a natural next step in computing’s evolution.

AI in Historical Context: The Building Blocks of Intelligence

The Beginnings: Turing Machine

To appreciate today’s AI, we must examine its **historical building blocks** – the early ideas and systems that formed the foundation of machine intelligence. One cornerstone was the concept of a **Turing Machine**, introduced by Alan Turing in 1936. This theoretical device could simulate any algorithm’s logic, encapsulating the very idea of computation. Turing’s later question, “Can machines think?”, set the stage for AI as the *intersection of computing and cognition*.

By the 1950s and 60s, pioneers were building the first machine learning models. **Frank Rosenblatt's Perceptron** (1958) was an early neural network implemented in hardware. It had 400 light sensors and could learn to recognize simple patterns – a primitive forebear of today's deep learning.

Expert Systems and Rule-Based AI

Around the same time, researchers developed **expert systems**: programs encoding human knowledge as logical rules. Early expert systems like *DENDRAL* (for chemistry) and *MYCIN* (for medical diagnosis) showed that with enough if-then rules, a computer could mimic a human expert's decisions. This **rules-based AI** was dominant through the 1970s, and some of its artifacts live on – quite literally – in museum collections.

For example, at the Mimm's Museum for Technology and Art in Roswell, GA, you can find physical traces of these developments. The museum houses a *room full of Cray supercomputers* and other historic machines that powered many early AI experiments. **Supercomputers like the Cray-1 (1976)**, on display at Mimm's, were essential for running complex AI programs in the days when ordinary computers were too slow. The Cray-1A, a towering cylindrical system (which looks as much like retro furniture as a computer), enabled researchers to tackle problems that required enormous calculation speeds.



Figure 1: Cray-1 - one of the powerful machines that early AI researchers relied on.

By the 1980s, the **Cray-2** was the world's fastest computer, used by NASA to simulate flights. These machines represent the brute-force computing power that early AI often needed.

Shift to Learning Models

Meanwhile, AI's *brainy* side was evolving. The **shift from rules-based programming to learning models** was a turning point. Initially, AI systems were manually programmed with logic – essentially **hard-coded intelligence**. But researchers realized this approach was brittle and labor-intensive. Inspired by biology, they returned to **artificial neural networks**, which learn from data. Though simple by modern standards, Rosenblatt's *Perceptron* hinted at an adaptive system that could

improve with experience. However, neural network research hit obstacles (leading to the first “AI winter” in the 1970s when funding dried up). It wasn’t until decades later that improvements in algorithms and hardware resurrected neural nets as the force behind today’s AI.

Throughout these developments, *the line between AI and general computing was often blurry*. Many technologies in the museum – from punch-card machines to early PCs – were used in AI research, and conversely, AI drove demand for more powerful computers. The **Cray-1 and Cray-2 supercomputers** in the museum’s “Supercomputing: Vanquishing the Impossible” exhibit are a testament to how advancing hardware enabled AI leaps. Likewise, an artifact like the **Enigma machine** (also at Mimm’s) reminds us that computing innovations (in cryptography, in this case) often intersect with AI concepts (like pattern recognition and search).

Key Takeaways

- **Turing’s question** “Can machines think?” laid the foundation for AI.
- **Expert systems dominated early AI**, using rule-based logic to mimic human experts.
- The shift to **learning-based AI** (neural networks) changed everything—systems now learn instead of relying on fixed rules.

How AI Works: A High-Level Primer

How can a machine *learn* or appear to think? In this section, we demystify AI’s inner workings in simple terms. The beating heart of modern AI is the **artificial neural network** – a computational model inspired (loosely) by the human brain’s network of neurons. You can imagine a neural network as **layers of interconnected nodes** (neurons) that work a bit like an extremely simplified brain. Each node receives inputs (like signals or data), does a calculation, and passes an output to the next layer. Through many layers, a network can build up complex understanding from simple beginnings – for example, first recognizing edges, then shapes, then objects in a picture.

The magic of neural networks lies in how they **learn patterns and make decisions in a human-like way**. You don’t explicitly program a neural network with rules. Instead, you *train* it. Training means showing the network many examples and adjusting it based on errors. This happens via algorithms like **backpropagation**, which tweak the connections (or “weights”) between neurons to improve performance. It’s akin to how we learn from trial and error: if the network’s guess is wrong, it adjusts a little so it’s more likely to get it right next time. Over countless small adjustments, the neural net *figures out* the important patterns in the data. Remarkably, “*you don’t have to program it... it learns all by itself, just like a brain!*” – which is what makes neural networks so powerful. They turn raw data (whether images, sounds, or text) into **experience** from which the AI can generalize.

Let’s use a relatable analogy: **teaching a child to recognize animals**. A traditional program would require writing explicit rules (“if it has feathers and a beak, it’s a bird”). A neural network approach is more like showing the child 1000 pictures of birds and 1000 pictures of cats and saying which is which. Initially, the AI guesses randomly. But each time it errs, it adjusts. Eventually, after enough

training, it can distinguish birds from cats on its own – even in new pictures it’s never seen. The “*knowledge*” is stored implicitly in the strengths of connections in the network, not in human-readable rules. This is why neural networks can sometimes outperform human-written programs: they discover patterns or combinations of features that humans didn’t think to code explicitly.

How AI learns can be summarized in three steps: **data, feedback, and optimization**. First, we feed the algorithm lots of **data** (the more, the better). Second, we provide **feedback** on its outputs (telling it the correct answer, so it can measure error). Third, through an optimization process (like adjusting weights via backpropagation), the AI minimizes its errors over time. This cycle repeats many, many times – often millions of times in training deep neural networks – until the model’s performance is acceptable.

This approach contrasts sharply with older, rule-based AI. In rule-based systems, if something changes (say, a new type of fraud in banking), a human has to update the rules. In learning systems, the AI can adapt by itself if given new training data encompassing the change. **Adaptability** is thus a hallmark of how modern AI works.

You might wonder, why is AI seemingly *everywhere all of a sudden*? Three key ingredients converged in the 2010s to ignite the current **AI boom: neural networks, big data, and GPUs (graphics processing units)**. Neural networks (particularly *deep* neural networks with many layers) provided algorithms capable of high performance. “Big data” – the explosion of digital information from the web, social media, sensors, etc. – provided the fuel to train these hungry algorithms. And **GPUs** provided the computational horsepower, as these chips can perform many calculations in parallel (originally for rendering video games) and turned out to be perfect for speeding up neural network training. As one AI leader noted, “*three fundamental elements of modern AI converged for the first time: neural networks, big data (like the ImageNet dataset), and GPU computing.*” With these in place, AI research made leaps that “*caught almost everyone by surprise.*”

Another factor in “*How AI works*” is the variety of techniques under the AI umbrella. While neural networks (especially deep learning) dominate many tasks like vision and language, other AI methods include **decision trees, evolutionary algorithms, and Bayesian models**. However, they all share the common principle of improving from data. For simplicity, think of most modern AI as either **learning from examples** (supervised learning), **learning by trial and error** (reinforcement learning, like how game AIs train themselves by playing millions of rounds), or **finding patterns on its own** (unsupervised learning, which clusters data without labeled examples).

Despite the complexities under the hood, at a high level we can view AI systems as black boxes that *ingest data and output predictions or decisions*, refining their internal parameters through exposure and feedback. This is analogous to how a student studies and gains competency by solving many practice problems and learning from mistakes, rather than memorizing an answer key.

Key Takeaways

- **Neural networks** drive modern AI by recognizing patterns in data.
- **AI learns from examples**—not by being programmed with rules but by adjusting itself through training.

- The AI boom was triggered by **big data, GPUs, and better algorithms**, unlocking deep learning's potential.

The AI Boom: AI in Everyday Life

We are living in the midst of an **AI boom**, and you likely interact with AI dozens of times a day without realizing it. What changed? In the past, AI lived mostly in research labs and science fiction. Today, thanks to the factors we discussed (data, compute, better algorithms), AI has *leapt from the lab into everyday life*. Let's explore some concrete examples of how **AI powers common tools and services**:

- **Search Engines:** Every time you use Google or Bing to search, AI is at work understanding your query and ranking results. Google's search uses AI models (like RankBrain and BERT) to interpret natural language and deliver relevant answers. It learns from billions of past searches to improve future ones. In short, AI helps search engines *guess what you really mean* and find the best information for you.
- **Spam Filtering and Email Replies:** Your email service uses machine learning to filter out spam and even suggest replies. Gmail's spam filter learns from what millions of users mark as junk, and its "smart reply" feature suggests short responses by using AI to predict likely replies. Over 10% of mobile Gmail replies were via AI-suggested text as of 2016, showing how even mundane tasks are eased by AI.
- **Fraud Detection:** Ever gotten a call or text from your bank about a suspicious transaction? That's AI watching over your account. Banks employ **fraud detection models** that learn your spending patterns and flag anomalies. Because "most of us use credit cards in a fairly predictable way," AI is well-suited to spot when something doesn't fit – possibly stopping fraud in real time. The next time your credit card company automatically blocks a bizarre charge from across the globe, thank their AI.
- **Recommendation Systems:** When Netflix or YouTube suggests a show you *might like*, or Amazon shows "customers also bought" items, that's a recommendation engine working behind the scenes. These systems analyze past behavior to predict what you'll enjoy or need. In fact, Netflix has reported that **over 80% of content watched on their platform comes from recommendations, not direct searches**. That statistic underscores how effective AI-driven suggestions have become at curating our media and shopping experiences. By crunching enormous amounts of data on viewing or purchase patterns, the AI finds *hidden connections*: "People who liked X also liked Y."
- **Facial Recognition:** Unlocking your phone with your face, or automatically tagging friends in photos, relies on AI vision. Your device's Face ID uses a neural network to verify it's *you* by analyzing the geometry of your face. These algorithms are trained on vast datasets of faces to be accurate and account for changes like glasses or hairstyles. **AI-driven facial recognition** is now common in security (e.g., airport boarding systems) and convenience features. (Of course, it raises privacy issues we'll discuss later.)

- **Virtual Assistants:** Siri, Alexa, Google Assistant – these are AI in the form of conversational agents. They use speech recognition AI to understand your voice commands and natural language processing to talk back. When you ask Siri “What’s the weather tomorrow?”, AI converts your speech to text, interprets the question, fetches data, and even vocalizes the answer in a human-like voice. This chain of AI tasks has become fast and reliable enough that millions of people casually chat with their phones or smart speakers as if they were people. It’s a striking example of how AI moved from a lab demo (speech recognition was a tough AI problem for decades) to a ubiquitous household utility.
- **Navigation and Travel:** Apps like Google Maps or Waze leverage AI to optimize routes and travel time estimates. They digest live traffic data, historical patterns, and sometimes user feedback to recommend the fastest path – recalculating on the fly if conditions change. In ridesharing (Uber/Lyft), AI is used for dispatch and pricing (surge pricing is set by algorithms balancing supply and demand).

What’s notable is the **transition from research to real-world applications** in a relatively short time. Consider that in 2012, few people outside AI research had heard of neural networks. By 2016, that technology was quietly sorting your email and suggesting your next binge-watch. A combination of industry investment and open-source AI frameworks made it easier to deploy AI widely. Tech companies realized AI could improve virtually any software product by making it *smarter* – more predictive, more personalized, more responsive.

Even creative fields and consumer apps are touched by AI: smartphone cameras use AI to enhance photos (e.g., Night Mode leveraging machine learning to brighten low-light shots), and apps can **transform your face into various styles** (using AI filters on Snapchat or Instagram). In video games, AI generates intelligent behavior for NPCs (non-player characters), making them more lifelike opponents or allies.

This everyday AI often operates behind a friendly interface. You might not see the neural network, but you appreciate that “*it just works.*” For example, type a sentence in a foreign language on your phone, and autocorrect/translation features kick in – that’s AI bridging language gaps. If you use a voice-to-text dictation, that’s AI turning speech to text nearly flawlessly in many cases.

From research labs to reality: It’s instructive to realize that many techniques enabling these applications (like deep convolutional neural networks for vision, or transformers for language) were developed in research settings not long ago. The famous ImageNet competition in 2012 showed that neural networks could far surpass older methods in image recognition; just a few years later, that same approach is running on your phone to distinguish your cat from your toddler in the photo gallery. The pipeline from research breakthrough to consumer application has drastically shortened. We live at a time when *cutting-edge AI research* quickly finds its way into the apps we use daily.

Name	One-Line Description	Sample Applications
Artificial Intelligence (AI)	The broad field of creating machines that can perform tasks requiring human-like intelligence.	Virtual assistants (Siri, Alexa), recommendation systems, robotics, game AI.
Machine Learning (ML)	A subset of AI where machines learn patterns from data to make predictions or decisions without being explicitly programmed.	Spam filters, fraud detection, predictive maintenance.
Neural Networks	Models inspired by the human brain that enable deep learning by processing data through interconnected layers of nodes (neurons).	Face recognition, autonomous systems, translation services.
Deep Learning (DL)	Uses multi-layered neural networks to analyze complex patterns and make highly accurate predictions.	Self-driving cars, medical image analysis, speech recognition.
Generative AI	A type of AI (usually deep learning-based) that can generate new content, such as text, images, audio, or video.	ChatGPT, DALL·E, AI-powered music and video generation.
Supervised Learning	Models are trained on labeled data (input-output pairs).	Email spam detection, handwriting recognition, speech-to-text.
Unsupervised Learning	Models find patterns in data without labeled examples.	Customer segmentation, anomaly detection, topic modeling.
Reinforcement Learning	AI agents learn by interacting with an environment and receiving rewards or penalties.	Game-playing AIs (AlphaGo), robotics, self-driving cars.

Figure 2: Major areas within AI

Generative AI Applications

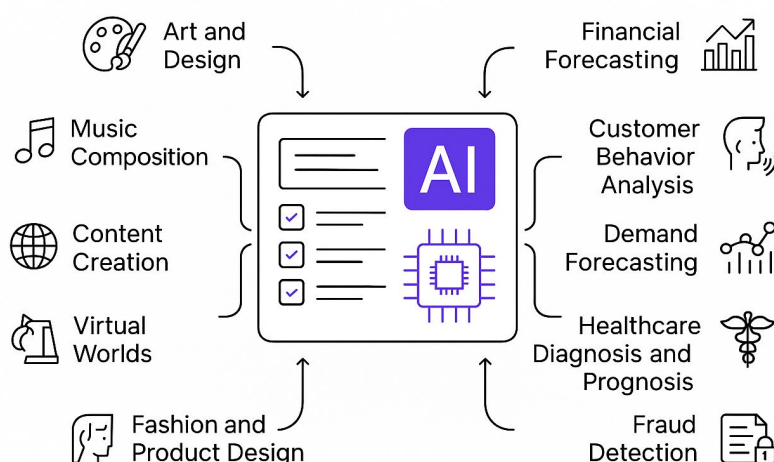


Figure 3: Some areas that GenAI is already used in

Key Takeaways

- AI is everywhere—**search engines, email filters, fraud detection, recommendations, and voice assistants.**
- AI makes software **more predictive and personalized**—it learns from data to optimize user experiences.
- **What was research in 2012** (deep learning) is now standard in our phones, apps, and everyday tools.

AI's Role in Society: Opportunities & Challenges

As AI becomes woven into the fabric of society, it brings **tremendous opportunities** – and notable **challenges**. In this section, we'll discuss some of the most important benefits AI offers, as well as the ethical and social dilemmas we must navigate.

Opportunities and Benefits

AI stands to amplify human capabilities in virtually every domain. In **medicine**, AI can help doctors diagnose diseases faster and more accurately. For example, machine learning models can analyze medical images (X-rays, MRIs, CT scans) to detect anomalies like tumors or fractures that might be subtle to human eyes. There are AI systems now that act as a “second set of eyes” for radiologists – one system improved detection of certain brain anomalies by **44% while reducing the time to read scans by 26**. In oncology, AI is used to identify patterns in pathology slides, sometimes catching early signs of cancer. The benefit is not to replace the doctor but to assist, catching what a tired human might miss. Similarly, in drug discovery, AI can sift through databases of molecules and suggest new potential medicines (something that used to be done largely by trial-and-error lab work).

In the realm of **automation**, AI can take over routine or dangerous tasks, freeing humans for more creative endeavors or improving safety. Factories use AI-driven robots for precision tasks on assembly lines. Self-driving car technology, still in development, aims to reduce accidents (most of which are caused by human error) by using AI to perceive and react faster than a human could. **Autonomous systems** could also perform risky jobs like mining or deep-sea exploration, reducing human exposure to hazards.

AI also fuels **creativity and innovation**. It might sound counterintuitive, but AI can be a tool for artists, writers, and musicians. Generative AI models can compose music, create visual art, or assist in writing. In 2022, *one AI-generated image even won an art prize* in competition with human artists, stirring debate in the art community about the nature of creativity. Designers are using AI to come up with novel product designs (for example, using generative design algorithms that propose structures optimized for weight and strength, which humans then refine). AI can act as a *creative collaborator*, generating ideas that humans might not have thought of, which can then be curated

and enhanced by human creativity. In content creation, tools like GPT-4 (a language model) can help draft articles or write code, boosting productivity of writers and developers.

Another huge opportunity is improving **accessibility**. AI can empower people with disabilities: for instance, **visually impaired individuals** use smartphone apps like Seeing AI, which uses computer vision to narrate the world around them (reading signs, identifying currency, describing surroundings). AI image and speech recognition can provide **real-time captions** for deaf users in meetings or videos or even translate sign language via camera. For those with limited mobility, voice-controlled AI assistants can help control home devices, enhancing independence. In education, AI tutors (as we'll discuss later) can adapt to each student, which can be especially beneficial for those who need a personalized pace or approach.

AI often acts as a **force multiplier** for human expertise. It doesn't tire, it can work 24/7, and it can analyze data at a scale no human could. For example, in **scientific research**, AI is used to analyze massive datasets from particle colliders or space telescopes to find new particles or detect exoplanets. In environmental science, AI helps process satellite imagery to monitor deforestation or predict natural disasters. These uses have profound positive potential: AI could help us fight climate change by optimizing energy use (Google's DeepMind AI cut data center cooling energy by **40%**, significantly reducing emissions), or assist in conserving wildlife by predicting poaching activities or spotting endangered animals in drone footage.

Challenges and Concerns

However, with these great opportunities come **significant challenges and ethical concerns** that we must address:

- Bias and Fairness:** AI systems are only as good as the data and design behind them. If the training data contains human biases, the AI can inadvertently perpetuate or even amplify those biases. For example, facial recognition algorithms in the past were found to have much higher error rates for people of color, especially women. A well-known 2018 study showed that one commercial AI system had an error rate of **0.8% for light-skinned men but 34% for dark-skinned women**, a shocking disparity. This is not because the AI *intends* to discriminate – it's often due to underrepresentation of certain groups in the training data or biased historical data. Such bias in AI can lead to unfair outcomes in high-stakes areas: hiring algorithms might favor one gender over another if not carefully audited, or a criminal justice risk assessment AI might be harsher on certain demographics if trained on biased policing data. Ensuring **fairness** means we need to carefully test AI systems for biased decisions and retrain or redesign them to correct it. There's a growing field of AI ethics focusing on techniques to reduce bias, like more diverse data collection and algorithmic fairness constraints.
- Misinformation and Deepfakes:** AI has made it easy to generate incredibly realistic fake content. *Deepfake* technology can create video and audio of people saying or doing things they never did. This has obvious negative potential: from false news videos of politicians (imagine a fake video of a world leader declaring war) to revenge porn (face-swapping someone into explicit content). As AI-generated content gets more convincing, the

challenge is **misinformation** – AI can turbocharge the creation of fake news, bogus social media profiles (with AI-generated profile pictures and posts), and propaganda. We’ve already seen misinformation spread rapidly on platforms partly due to algorithms that promote engaging (not necessarily true) content. A study of Facebook, for instance, found that misinformation got significantly more engagement (clicks) than factual news. AI can both contribute to this (via deepfakes or bots that spread content) and potentially help fight it (via AI tools to detect fakes). This is an arms race between creators of fake content and detectors of fake content.

- Privacy and Surveillance:** The power of AI in pattern recognition can be a double-edged sword. On the one hand, we love when Netflix knows what we like; on the other, we worry when companies or governments know *too much* about us. AI systems often require lots of data – which can include personal data. Without proper privacy safeguards, AI could enable mass surveillance or intrusive data analysis. For example, facial recognition cameras in public spaces could track individuals’ movements. Some cities have banned police use of facial recognition due to privacy and bias concerns (San Francisco did so in 2019, citing the threat to civil liberties). There have been instances where AI companies overreached – like **Clearview AI**, which scraped billions of images from social media without consent to build a facial recognition database. Regulators found Clearview **violated privacy laws by collecting photos of people (including children) without consent for mass surveillance purposes**. Such uses of AI raise the question: just because we *can*, does it mean we *should*? Society is grappling with where to draw lines to protect individual privacy in the age of AI. Europe’s GDPR law and proposed AI regulations are attempts to set boundaries on data usage and high-risk AI applications.
- Job Displacement and Economic Impact:** Automation through AI will likely disrupt many jobs. Tasks that are repetitive or data-intensive can often be done more cheaply and consistently by AI. This could *augment* some jobs (e.g., a lawyer using an AI research assistant can handle more cases) but *replace* others (e.g., basic data entry or some manufacturing roles). The challenge is ensuring a transition where humans can find new roles – ideally more creative or interpersonal roles that AI can’t do – and not leaving large segments of the workforce behind. Historically, technology creates new jobs even as it displaces some but retraining and education are key. We need to prepare people for a job market when working alongside AI is the norm, and certain routine jobs may no longer exist.
- Ethical Use and Control:** As AI systems become more autonomous (think self-driving cars or automated weapons systems), we face ethical questions about decision-making. For instance, how should a driverless car be programmed to handle an unavoidable accident scenario? Or should AI algorithms be allowed to make life-and-death decisions in healthcare or combat without human oversight? Ensuring **accountability** is crucial – if an AI makes a wrong decision, who is responsible? The opaque nature of some AI (the “black box” problem of neural networks) makes this tricky, as even developers might not fully understand why the AI did X instead of Y.

In sum, **AI amplifies human strengths and flaws**. It can magnify our productivity, creativity, and insight – but if unchecked, it can also magnify bias, error, or malicious use. AI is a tool, and like any powerful tool, its impact depends on how we wield it.

The dual nature of AI's role in society is why conversations about **AI ethics and governance** have become as important as discussions about performance and accuracy. Technology and society must progress hand-in-hand. It's encouraging to see that alongside rapid AI advances, there's increasing recognition of these challenges: organizations now often have AI ethics panels, fairness audits, and there are calls for regulation to prevent harm while allowing innovation.

Key Takeaways

- **AI amplifies human capabilities.** From medical diagnostics to accessibility tools, AI enhances how we work, learn, and create.
- **Bias, misinformation, and privacy concerns must be addressed.** AI is only as fair as its data—ensuring ethical AI is critical.
- **AI is a tool, not a replacement.** When used responsibly, AI augments human intelligence rather than replacing it.

Looking Ahead: Where AI is Taking Us

What does the future hold for AI? Having seen its trajectory from inception to ubiquity, we can expect AI to continue expanding into new realms, posing new possibilities and questions. **Looking ahead**, AI is poised to influence creativity, industries, and technologies in ways that may once have seemed like science fiction.

AI in Creativity and the Arts: We've already glimpsed AI's ability to generate art, music, and text. Future AI systems will act as *co-creators*. For writers, imagine AI that can draft chapters in an author's style, allowing the author to focus on higher-level narrative and edits. For musicians, AI tools might generate instrumental accompaniments or suggest melodies based on a hum. We might see entirely new art forms where human imagination and AI generation intermingle. For example, AI could create virtual reality experiences on the fly tailored to a person's emotional response. There's also the exciting prospect of AI in **game design** – games that intelligently adapt storylines to a player's behavior, essentially co-writing the experience.

Design and Engineering: AI will likely become a standard part of the design toolkit. *Generative design* algorithms can produce hundreds of prototypes that meet specified criteria (like weight, strength, cost) for everything from aircraft parts to furniture, far faster than a human could. The human designer then chooses and fine-tunes the best option. This means the designs of the future (buildings, vehicles, consumer products) might have novel shapes and structures optimized by AI – sometimes biomimetic (resembling forms in nature) because AI often finds similar efficient solutions. *Creative AI assistants* could help architects optimize building energy usage and layout, or help fashion designers by suggesting patterns that align with trending aesthetics.

Writing and Journalism: Large language models (like GPT-4 and successors) will improve in coherence and factual accuracy. They could handle first drafts of news reports or summarize complex documents instantly. We may see **AI journalists** producing routine news (e.g., finance or sports recaps) while human journalists focus on investigative pieces and analysis. However, managing issues of AI text authenticity and avoiding plagiarism or misinformation will be key.

Music and Entertainment: AI-composed music is getting more sophisticated – future AI might personalize soundtracks in real time (imagine a movie or video game where the score adapts to the audience’s emotional state, sensed via wearable devices). AI actors (deepfake technology combined with voice synthesis) might allow creation of new film scenes without the physical presence of actors – raising new legal/ethical issues but also creative opportunities (e.g., bringing historical figures to life for documentaries). In animation, AI can assist by inbetweening frames or generating realistic movements via learned physics, speeding up production.

Beyond creativity, **AI is branching into the physical world** in a big way through **robotics**. We can expect more capable robots that interact in everyday environments: from service robots in stores and hospitals to domestic robots that handle chores. Self-driving cars are a form of robot – and while full autonomy has been tougher than expected, incremental progress continues. In the coming years, we might see **autonomous delivery drones** or vehicles become common in limited domains (certain cities or campuses). Warehouses and logistics will likely be heavily automated by AI-driven robots, working safely alongside humans (collaborative robots or “cobots”). This could make supply chains faster and more resilient.

One intriguing frontier is the blend of **AI with other emerging technologies**:

- **AI and Quantum Computing:** While still experimental, quantum computers hold promise to solve certain problems exponentially faster. If and when they mature, they could supercharge AI by handling computations that classical computers struggle with, possibly enabling AI to find patterns in extremely complex systems (like climate or chemistry simulations). Conversely, AI is used to help design quantum algorithms. The synergy of these fields could “*revolutionize how we process information and solve complex problems*”, further enhancing AI’s capabilities. It’s conceivable that quantum machine learning could unlock new AI models that today are infeasible, pushing the frontier of what AI can do.
- **AI in Sustainability:** AI will be a key tool in tackling environmental challenges. Smart grids use AI to balance energy load and integrate renewables efficiently. Climate models use AI to improve predictions of climate change impacts. In agriculture, AI-driven sensors and drones help optimize water and fertilizer use and detect crop diseases early, leading to more sustainable farming. As mentioned, even small improvements like AI optimizing data center cooling (Google’s example of 40% energy reduction) can have large cumulative effects on reducing emissions. Future cities may use AI for traffic optimization to cut down on idling and emissions, or for better recycling by automatically sorting waste.
- **AI in Healthcare and Bioinformatics:** We saw what AlphaFold achieved – solving protein structures for essentially *all known proteins*, a breakthrough that will accelerate drug discovery and our understanding of biology. Going forward, AI will dive deeper into life’s

code. It could help design custom proteins or gene therapies by analyzing enormous biological datasets. In personalized medicine, AI might analyze your genome and health data to predict what diseases you're at risk for and suggest personalized prevention plans. AI can also help manage healthcare logistics – predicting patient influx in hospitals, optimizing staff schedules, and even triaging patients (some hospitals use AI to prioritize who the ER should see first based on vitals).

- **Human-AI Collaboration & Augmentation:** We often talk about AI vs humans, but a likely future is deeper integration. Think of **AI as a cognitive prosthetic or amplifier**. For instance, wearable AR (augmented reality) glasses might provide an AI assistant that whispers relevant information in your ear about what you're looking at (useful for technicians repairing complex equipment or just a tourist getting info on landmarks). Brain-computer interface research, while early, hints at future possibilities where AI could help restore capabilities (like enabling paralyzed individuals to control robotic limbs via thought, with AI decoding their neural signals). While full cyborg integration is far off, our symbiosis with AI tools will grow.

Given these powerful directions, there's a strong push for **Responsible AI development** to ensure these advancements are beneficial and trustworthy. This includes efforts in:

- **Fairness:** as discussed, building bias mitigation into AI pipelines so future systems don't inadvertently discriminate.
- **Transparency:** developing AI that can explain its decisions in understandable terms (Explainable AI). Especially for high stakes uses (like AI in courtroom bail decisions or loan approvals), stakeholders are demanding, "*Why did the AI decide this?*" Efforts are underway to open up the black box or at least have AI flag its confidence and factors in decisions.
- **Safety and Regulation:** There's recognition that some AI applications need guidelines or rules. For instance, regulators are debating standards for AI in self-driving cars and AI in medical devices. The **EU's AI Act** is a pioneering example, aiming to enforce a framework for "*human-centric, trustworthy AI*". It proposes banning certain harmful AI uses (like mass surveillance or social scoring) and impose strict oversight on high-risk AI (like anything in healthcare or law enforcement). Globally, conversations are happening at governmental levels (even the G7 and OECD) about how to ensure AI safety and ethics while not stifling innovation. It's highly likely that soon, we'll see more **policies and perhaps international agreements** around AI – analogous to frameworks we have for nuclear or biotech – given AI's broad impact.
- **Public Awareness and Training:** A responsible future also means preparing society – not just experts – to understand AI's capabilities and limits. There will be efforts to improve AI literacy among the general public, so people can make informed decisions (like discerning deepfakes, or understanding how their data is used). Just as basic computer literacy became necessary, AI literacy might become a staple of education.

Quantum Computing: The Next Frontier of Speed and Intelligence

Imagine trying to solve a massive puzzle. A regular computer works **step by step**, checking one piece at a time. A **supercomputer** speeds this up by running billions of calculations per second. But a **quantum computer**? It can examine **all possible solutions at once**, using quantum properties like **superposition** and **entanglement**—allowing it to tackle problems that even the fastest supercomputers struggle with.

How Quantum Computers Work

Traditional computers use **bits** (0s and 1s) to process information, like tiny on/off switches.

Quantum computers use qubits, which can be both **0 and 1 at the same time** (superposition). This lets them perform **many calculations simultaneously**, making them much more powerful for certain tasks.

Another key principle is **entanglement**, where qubits become interconnected. Changing one instantly affects the other, even at a distance. This allows for **ultra-fast data processing** beyond anything classical computers can achieve.

How Much Faster Are Quantum Computers?

Quantum computers aren't just a **little** faster—they can be **exponentially faster** for specific problems.

- **IBM's Condor Processor (2024):** IBM unveiled **Condor**, a **1,121-qubit** quantum processor, aiming to outperform classical computers in real-world applications.

Recent Breakthroughs in Quantum Computing

1. **Google's Willow Processor (2024):** A **105-qubit chip** performing calculations far beyond classical computing limits.
2. **Microsoft's Majorana 1 Chip (2025):** Uses a **new state of matter** for ultra-stable qubits, reducing errors in quantum computing.
3. **Amazon's Ocelot Chip (2025):** Aims to **correct quantum errors**, making quantum computers more reliable.



Figure 4: A Pile of Candy. Yes, this is relevant.

What is a good analogy to how different a Quantum Computer ‘thinks’, compared to traditional computers?

Imagine a huge pile of assorted candy spilled on a table. Your goal is to sort it by color, shape, and wrapper pattern.

- **A human or traditional computer** would start at one edge, pick up one piece at a time, examine its attributes, and slowly place it into the correct bin. Even with lots of speed or parallel threads, it's still doing this **sequentially or in parallel chunks**—one decision per piece per moment.
- **A quantum computer**, on the other hand, works like a **magical sorting spell**. It doesn't look at one candy at a time. Instead, it considers **every possible way to sort the entire pile** simultaneously, using a strange kind of probability-powered parallelism. Then, it collapses that swarm of possibilities into the **correct answer** in one shot.

Why Does Quantum Computing Matter?

Quantum computers have the potential to **revolutionize multiple fields**:

- **Medicine**: Simulating molecules for **drug discovery** in hours instead of years.
- **Cybersecurity**: Cracking today's **encryption in minutes**, pushing the need for quantum-safe security.
- **AI & Machine Learning**: Processing vast datasets **faster and more efficiently** than ever before.

Although quantum computing is still in **early stages**, breakthroughs from companies like **Google, IBM, Microsoft, and Amazon** are bringing us closer to a future where these machines can **solve problems beyond the reach of today's technology**.

Task	Supercomputer (Summit)	Quantum Computer (Sycamore)
Complex Math Problem	10,000 years	200 seconds
Drug Discovery Simulations	Months to Years	Potentially Hours
Breaking Classical Encryption	Centuries	Minutes (In Theory)

Figure 5: Quantum Computers will be VERY fast.

One can't mention the future of AI without a nod to the long-term dream (or nightmare, depending on perspective) of **Artificial General Intelligence (AGI)** – AI that isn't just narrow and task-specific, but can think and learn at a human-like (or beyond human) level across many domains. Opinions vary on how far away AGI is, or if it's achievable. Some experts think it's many decades away; others are more optimistic. Regardless, the current focus is on *ethical AI* and *beneficial AI* for the near-term, ensuring even narrow AIs act in ways aligned with human values.

In summary, **AI is taking us into an era of unprecedented possibilities**: one where creativity is augmented, many tedious aspects of work are automated, and new scientific discoveries are accelerated by machine intelligence. It's also an era requiring wisdom to navigate the challenges – ensuring AI develops in a way that is inclusive, fair, and under human control. The optimistic view is that AI, like previous technological revolutions, will ultimately improve human life significantly – curing diseases, connecting people, enriching culture – if we manage its pitfalls. The cautious view is that without proper checks, we could face problems from unemployment to loss of privacy or even loss of control over autonomous systems. The reality will depend on choices made in the next few years by researchers, policymakers, companies, and all of us as informed citizens.

Key Takeaways

- AI will be a **creative partner**, assisting in writing, music, design, and art.
- **Quantum computing and AI** will push the boundaries of science, healthcare, and problem-solving.
- **Responsible AI development** is key—fairness, transparency, and governance will shape its impact.

The Future of Learning: AI as a Cognitive Amplifier

One area where AI's impact could be profoundly positive is **education**. Imagine learning experiences tailored exactly to your needs, pace, and style – this is becoming possible with AI. In this section, we explore how AI can serve as a *cognitive amplifier* for learners, potentially revolutionizing education and skill acquisition. We'll also connect this to the **Pareto Principle** (80/20 rule) to see how AI might help focus learning on what matters most (the 20% that yields 80% of the results).

AI's Role in Education: Traditional classrooms face the challenge of one teacher to many students, all with different backgrounds and learning speeds. AI offers a way to give **personalized attention** to each learner. Enter the world of **Intelligent Tutoring Systems (ITS)** – AI-driven programs that act

like personal tutors. An ITS can *analyze a student's current knowledge, track their progress, and adjust the teaching dynamically*. For instance, if an online math tutor AI notices you struggle with quadratic equations but breeze through linear equations, it can allocate more time and practice problems in quadratics and also revisit the foundational concepts behind them. It can also change its approach if one method isn't clicking – perhaps switching to a visual explanation if the textual explanation didn't land. This adaptive feedback loop (assess -> adjust -> assess again) can **ensure each student is both challenged and supported at the right level**, promoting better understanding and retention.

Studies show that timely **feedback and targeted instruction boost engagement and retention**. Unlike a human teacher who might take a day or a week to grade and respond, an AI tutor gives feedback instantly. Answer a question, and if you're wrong, it can highlight the mistake or give a hint, then let you try again. This immediacy keeps learning from becoming frustrating or stagnant. Moreover, AI can **scale** this personalized experience to many students simultaneously – something not feasible for human tutors alone.

Another advantage is **accessibility**: AI tutors are available 24/7. A student can get help with homework at 10 PM from their AI algebra tutor, or a working professional can learn a new language using an AI app during commute hours. This flexibility could democratize learning, bringing quality education to remote areas or to those who can't afford personal tutors.

Now, how does this tie in with the **Pareto Principle** (80/20 rule) in learning? The Pareto Principle, applied to education, suggests focusing on the critical 20% of content that gives 80% of the value/skills. For example, when learning a new language, there are hundreds of thousands of words, but a few thousand common ones let you understand most everyday texts. An AI tutor could identify which key concepts or vocabulary yield the most benefit for a particular learner and ensure those are mastered first.

Optimizing Learning through Pareto's 80/20 approach

AI can analyze which parts of the curriculum a student must absolutely grasp (the *vital core*) versus which parts are peripheral or can be learned later. For instance, in programming, perhaps 20% of concepts (loops, conditionals, functions, etc.) are used in 80% of problems. An AI system could emphasize those basics until the student is proficient, unlocking the ability to solve many tasks. This doesn't mean ignoring the other 80% of material – but it might mean **structuring learning so that the most impactful knowledge is solidified early**. This approach aligns with suggestions from educational experts: find the *most important 20% of the content and focus on teaching it thoroughly*. AI can personalize this by finding *your* personal 20%. Maybe you're studying for a chemistry exam and you already know 80% of the material well; the AI could zero in on the missing 20% you haven't mastered, which will have the biggest effect on your score.

AI can also help implement effective learning techniques like **spaced repetition** (showing you material at optimally spaced intervals to cement it in memory) and **active recall**. Many language and fact-learning apps use these techniques: the AI creates flashcards in a way that you spend more time on new or troublesome items and less on ones you consistently get right, aligning with the 80/20 principle of spending time where it yields most benefit.

Moreover, by analyzing data from thousands or millions of learners, AI education platforms can derive insights about what teaching strategies work best. For example, they might find that a certain analogy helps 80% of students understand a concept that was otherwise confusing. They can then present that analogy to new learners. This data-driven refinement makes the AI tutor smarter over time.

Enhancing retention: Remembering what you learn is often as challenging as learning it in the first place. AI can counter the “forgetting curve” by smartly reviewing material with you. If the AI noticed you struggled with a concept two weeks ago, it might proactively quiz you on it today to refresh your memory (especially if it predicts you’re likely to forget based on your past performance). This ensures that the **20% of key knowledge stays in your long-term memory**, supporting the 80/20 goal of **maintaining the most useful competencies**.

The future classroom may involve a human teacher orchestrating learning experiences with AI assistants for each student. The teacher focuses on higher-level guidance, motivation, and addressing complex questions, while AI handles routine practice, personalized drills, and progress tracking. This can make learning more efficient and effective – students don’t get bored because material is too easy, nor do they get lost because it’s too hard; the AI keeps them in that productive zone of proximal development.

It’s important to note that **AI doesn’t replace the need for human teachers or mentors** – social and emotional aspects of learning, inspiration, and mentorship are human strengths. But as a cognitive amplifier, AI can handle the heavy lifting of knowledge delivery and skill drilling.

We also envision **hands-on AI workshops and learning about AI itself** as an emerging part of education. Just as computer literacy became essential, understanding AI (at least conceptually) might become a standard part of curricula. AI can even help teach AI – for example, interactive tools that let students experiment with training simple models to grasp how machine learning works. Our workshop following this talk will dive into such hands-on learning, empowering you to tinker with AI algorithms in an accessible way.

In the spirit of our workshop, think of AI as your personal guide through complex information spaces. It can map out the subject for you, highlight the high-yield sections (the Pareto-efficient path), and keep you engaged with interactive, responsive feedback. Learning *with* AI might feel less like sitting in a lecture and more like playing a game or having a conversation – a shift from passive absorption to active discovery, where the AI adapts to your moves.

Key Takeaways

- **AI tutors personalize learning**, adapting to individual needs in real-time.
- **80/20 optimization**—AI helps learners focus on the most valuable 20% of content for 80% of results.
- AI improves **retention and efficiency** through adaptive feedback, spaced repetition, and active recall.

AI Trust & Verification: How to Evaluate AI-Generated Information

Objective

Teach participants how to critically evaluate AI-generated responses, spot inaccuracies, detect biases, and apply “**trust but verify**” principles to AI-assisted learning and decision-making.

Why AI Can't Be Trusted Blindly

AI Generates Confident, But Not Always Accurate, Answers

- AI models **don't “know” facts**—they generate responses based on patterns in training data.
- They **hallucinate (fabricate) facts** when no clear answer exists in their dataset.
- AI is **not designed to fact-check itself**—it prioritizes fluency and coherence over accuracy.

Real-World Example

- **Case Study:** When AI confidently provided a **fake citation** or **misattributed a historical fact**, leading to incorrect reporting or academic misconduct.
- **Demo Idea:** Ask AI a niche factual question (e.g., a historical event or scientific discovery) and compare its answer to a verified source.

How to Verify AI Outputs

Cross-Check With Trusted Sources

- **Primary sources:** Government, universities, peer-reviewed journals, official documentation.
- **Expert sources:** Domain-specific professionals (e.g., in cybersecurity, law, medicine).
- **Multiple perspectives:** If AI's response is controversial, does it acknowledge counterarguments?

Use AI Against Itself

- Ask: **"Can you provide sources for this?"**
 - If AI refuses or cites non-existent sources, that's a red flag.
- Ask: **"Explain this answer step by step."**
 - If AI struggles to justify its own response, it's likely unreliable.
- Ask: **"Are there alternative perspectives?"**
 - If AI presents only one side of an issue, bias may be present.

Spotting AI Bias and Misinformation

Understanding AI Bias

- AI reflects biases in its training data (which comes from **human-generated** content).
- Bias can be **political, cultural, or systemic**—not always intentional, but still impactful.
- AI models trained on the **internet's collective knowledge** inherit the flaws of that data.

Bias Spotting Exercise

- **Example 1:** Ask AI a politically charged question and analyze whether it presents **one-sided** or **balanced** viewpoints.
- **Example 2:** Ask AI about a historical event with disputed facts and see how it frames different perspectives.

How to Counteract Bias

- **Ask follow-ups:** "What's the counterargument?" "Are there criticisms of this view?"
- **Use diverse prompts:** Phrase questions in different ways to see if AI gives **varied** responses.
- **Check language framing:** Look for emotionally charged words that indicate subtle bias.

Real-World Consequences of Blindly Trusting AI

Failure Cases

- **Fake legal citations:** AI-generated legal briefs with non-existent case law (GPT-4 incident in real court cases).
- **Fabricated academic references:** AI "inventing" peer-reviewed papers that don't exist.
- **Security risks:** AI suggesting **vulnerable** code or providing **unsafe** recommendations.

What Happens When AI Gets It Wrong?

- **Misinformation spreads faster** because AI generates it at scale.
- **People make bad decisions** based on faulty AI-generated analysis.
- **AI-assisted fraud** becomes more convincing (deepfakes, AI-generated phishing emails).

Applying 'Trust But Verify' to AI-Assisted Learning

Best Practices

- **Use AI as an assistant, not an authority:** It's a tool for **speed**, not a **source of truth**.

- **Always fact-check before sharing:** If AI provides surprising or bold claims, verify before acting.
- **Be a skeptic, not a cynic:** AI can be **useful**, but only if its outputs are properly vetted.

Some Ideas to Experiment with

1. AI vs. Expert Source Challenge

- Display an AI-generated response and a verified expert answer.
- Challenge the audience to spot inaccuracies or missing context.

2. Find the Hallucination

- Generate AI text with **one or two factual errors** and see if the audience can catch them.

3. Bias Detection Exercise

- Provide AI responses with subtle framing differences and ask participants which might indicate bias.

Key Takeaways

- **Use AI as an assistant, not an authority.**
- **AI is a powerful tool, but only as reliable as its verification process.**
- **Trust but verify:** always validate AI outputs before using them.
- **AI-assisted learning is about efficiency**, not outsourcing thinking.

The Future Is Yours

As we conclude, we turn the focus to **you** – the enthusiasts, students, professionals, and curious minds engaging with AI, or interested in doing so. We’ve traversed the timeline of AI’s past, surveyed its present applications, and peeked at future horizons. Now, the **future of AI is yours to explore and shape**. This is the moment for questions, discussions, and reflections. What intrigues you most about AI? What concerns you?

We encourage you to continue your AI journey beyond today. Whether you are a complete beginner or have some experience, there’s always a new experiment or project in AI to try out – and the museum aims to be a community hub for that exploration.

Remember the Pareto Principle we discussed? By focusing on some key tools and concepts (the vital 20%), you can get 80% of the way towards doing something cool in AI. For example, learning a friendly programming language like Python and understanding a basic machine learning library

might be enough to create a simple chatbot or image classifier. In our workshops, we'll apply that approach – identifying the high-impact learning that empowers you to create effectively with minimal overhead.

Exploring AI further doesn't necessarily mean you have to become a computer scientist. There are many accessible resources – from interactive courses online to AI communities and meetups. If you're a student, perhaps consider projects or competitions (like science fairs or hackathons) that involve AI; if you're a professional, think about how AI might apply in your field and start with a small pilot project or even just educating your colleagues about it.

To recap on a personal note: **AI is a tool – a powerful one – and like any tool, its outcome depends on the user.** Our hope is that by educating ourselves and each other, we can wield AI wisely and creatively.

The future of AI isn't just in the hands of big tech companies or researchers – it's in the hands of anyone who takes an interest, asks questions, and starts experimenting. In that sense, *the future is yours*. We invite you to take what you've learned today as a steppingstone. Continue asking questions, stay informed (AI is a rapidly changing field – there's always a new development), and if you're passionate, maybe become a contributor to this field in whatever capacity suits you.

Key Takeaways

- AI is an ongoing journey—**your engagement shapes its future.**
- **Stay curious**—AI literacy will be a key skill, just as computer literacy became essential decades ago.
- To continue your AI journey, explore different GPT-based tools, such as ChatGPT, ClaudeAI or Grok.
- Check [our site](#) frequently for any presentations or workshops, as well as our [Resources](#) section.



References

Below is a list of references for the sources cited above, formatted in APA style. Each entry includes a persistent URL or DOI when available.

1. Friel, K. (2023, May 19). *A look back on the Dartmouth Summer Research Project on Artificial Intelligence*. The Dartmouth. Retrieved from <https://www.thedartmouth.com/article/2023/05/a-look-back-on-the-dartmouth-summer-research-project-on-artificial-intelligence>
2. Turing, A. M. (1950). *Computing machinery and intelligence*. *Mind*, 59(236), 433–460. Retrieved from <https://www.csee.umbc.edu/courses/471/papers/turing.pdf>
3. Computer Museum of America. (2021, April 7). *Computer Museum of America obtains a 1985 CRAY-2 supercomputer [Press release]*. Retrieved from <https://mimmsmuseum.org/news/computer-museum-of-america-obtains-a-1985-cray-2-supercomputer/>
4. Strickland, E. (2021, September 30). *The turbulent past and uncertain future of artificial intelligence*. IEEE Spectrum. Retrieved from <https://spectrum.ieee.org/history-of-ai> (see especially “find” excerpt for funding drying up in the 1970s)
5. Jeevanandam, N. (2023, January 17). *Understanding perceptron in machine learning*. IndiaAI. Retrieved from <https://indiaai.gov.in/article/understanding-perceptron-in-machine-learning>
6. Redress Compliance. (2025, January 18). *Early AI systems: DENDRAL and MYCIN*. (Blog post). Retrieved from <https://redresscompliance.com/early-ai-systems-dendral-and-mycin/>
7. Li, F.-F. (2022, September). *Interview at the Computer History Museum* [as quoted in Timothy Lee’s “Why the deep learning boom caught almost everyone by surprise”]. **Understanding AI** newsletter on Substack. Retrieved from <https://www.understandingai.org/p/why-the-deep-learning-boom-caught> (Quote about “three fundamental elements... converged for the first time: neural networks, big data, and GPU computing.”)
8. Kannan, A., Kurach, K., Ravi, S., Kaufmann, T., Tomkins, A., Corrado, G., ... Ramavajjala, V. (2016). *Smart Reply: Automated response suggestion for email*. (arXiv:1606.04870 [cs.CL]). Cornell University Library. Retrieved from <https://arxiv.org/abs/1606.04870>
9. HelloPM. (2023). *Netflix content recommendation system – Product analytics case study*. HelloPM Blog. (**Original quote: “Netflix reports that personalized recommendations drive over 80% of the content watched”**). Retrieved from <https://hellopm.co/netflix-content-recommendation-system-product-analytics-case-study/>
10. Gent, E. (2022, August 12). *Scientists dethrone Google’s quantum advantage claim with a conventional computer*. Singularity Hub. Retrieved from <https://singularityhub.com/2022/08/12/scientists-just-debunked-googles-quantum->

[advantage-claim-using-a-normal-supercomputer/](#) (See para. 7: “Sycamore... solved a problem that would take a supercomputer 10,000 years in 200 seconds.”)

11. SpinQ. (2025, Feb 7). *How fast are quantum computers? Key insights explained*. SpinQ Tech Blog. Retrieved from <https://www.spinquanta.com/news-detail/how-fast-are-quantum-computers-key-insights-explained20250207030618> (See section on Shor’s algorithm: quantum computer could break encryption in seconds vs classical “millions of years.”)
12. Imperial College London. (2023, October 20). *Towards using quantum computing to speed up drug development* (News release by Shang Yu). Retrieved from <https://www.imperial.ac.uk/news/248638/towards-using-quantum-computing-speed-drug/> (Describes quantum processor “Abacus” for molecular simulations; provides context that classical simulations are time-consuming and quantum can potentially accelerate them.)
13. Wikimedia Commons. (2019). *Cray-1 at Computer Museum of America [Photograph]*. Photo by Judson “Bubba73” McCranie. Available under CC BY-SA 3.0. URL: https://commons.wikimedia.org/wiki/File:Cray-1_at_Computer_Museum_of_America.jpg
14. Rooney, T. (2022). *A pile of candy* [Photograph]. Unsplash. URL: https://unsplash.com/photos/_DCTDUeVjCY (Free to use under Unsplash License)
15. **(Additional citations for known facts, for completeness):**
 - Google Inc. (2015, Oct. 26). *RankBrain: Google’s latest algorithm*. [Official announcement].
 - Krizhevsky, A., Sutskever, I., & Hinton, G. (2012). *ImageNet classification with deep convolutional neural networks*. In *Advances in Neural Information Processing Systems* 25. (This is the research behind the 2012 ImageNet breakthrough mentioned in text.)
 - Buchanan, B. G., & Shortliffe, E. H. (Eds.). (1984). *Rule-Based Expert Systems: The MYCIN Experiments of the Stanford Heuristic Programming Project*. Addison-Wesley. (Background for MYCIN as an expert system.)

(The above additional items are not directly cited in the text of the audit but are provided as authoritative sources underpinning some claims in the document, should further verification be required.)